


Recurrent Neural Networks (RNN)



Saeed Sharifian

Recurrent Neural Networks

- ▶ The idea behind Recurrent neural networks (RNN) is to make use of sequential data.
 - Until here, we assume that all inputs (and outputs) are independent of each other.
 - It is a bad idea for many tasks, e.g., predicting the next word in a sentence (it's better to know which words came before it).
- ▶ They can analyze time series data and predict the future.
- ▶ They can work on sequences of arbitrary lengths, rather than on fixed-sized inputs.
- ▶ Neurons in an RNN have connections pointing backward.
- ▶ RNNs have memory, which captures information about what has been calculated so far.



Recurrent Neural Networks

We can process a sequence of vectors \mathbf{x} by applying a **recurrence formula** at every time step:


$$h_t = f_W(h_{t-1}, x_t)$$

new state

some function
with parameters W

old state


input vector at
some time step



Notice: the same function and the same set of parameters are used at every time step.

Recurrent Neural Networks

The state consists of a single “*hidden*” vector \mathbf{h} :



$$h_t = f_W(h_{t-1}, x_t)$$




$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

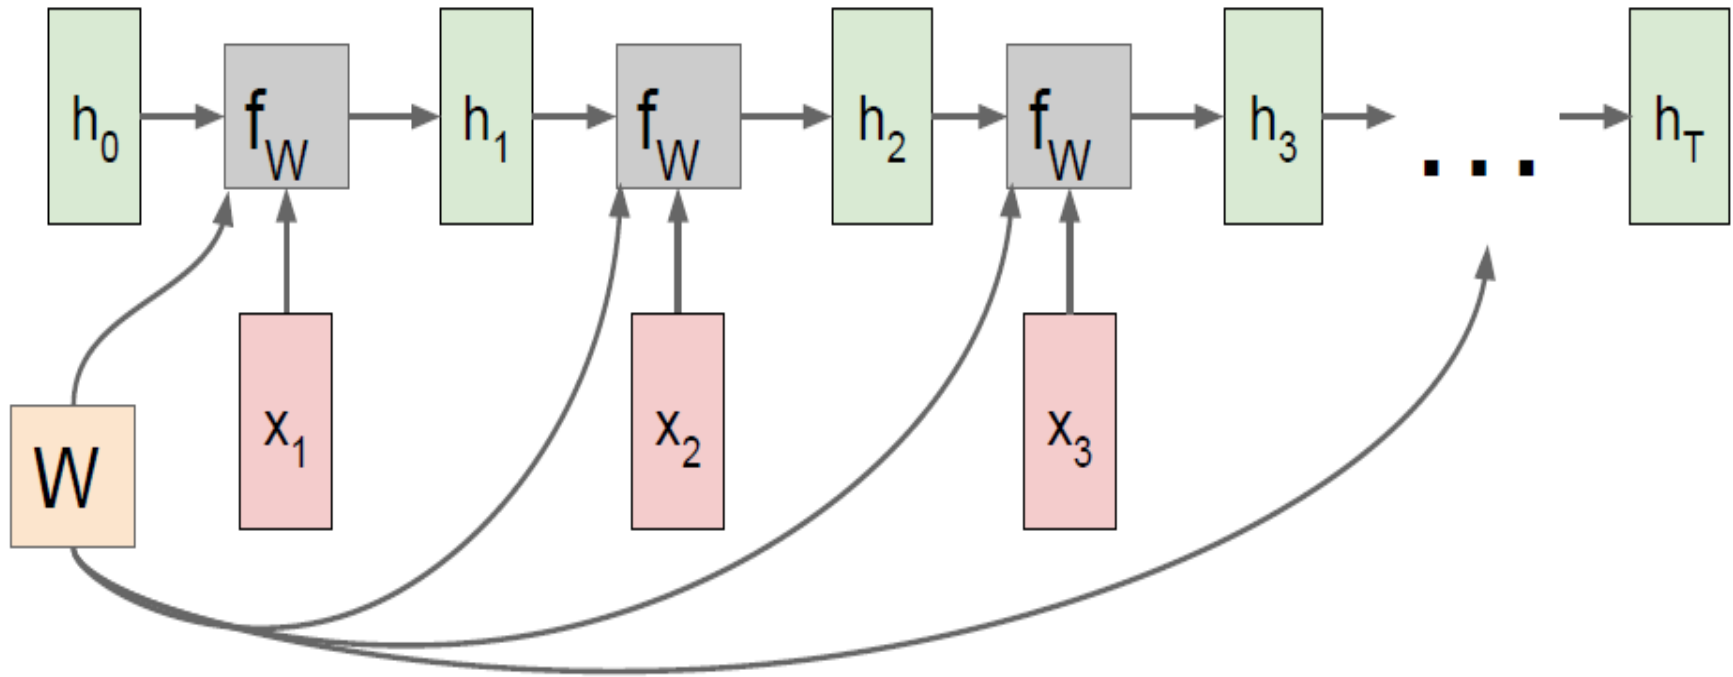
$$y_t = W_{hy}h_t$$

Recurrent Neural Networks

- **Unfolding the network:** represent a network against the time axis.
 - We write out the network for the **complete sequence**.
- For example, if the sequence we care about is a **sentence of three words**, the network would be **unfolded into a 3-layer** neural network.
 - One layer for **each word**.




Recurrent Neural Networks



Re-use the same weight matrix at every time-step


Recurrent Neural Networks

- ▶ $h^{(t)} = f(u^T x^{(t)} + wh^{(t-1)})$, where f is an activation function, e.g., **tanh** or **ReLU**.
- ▶ $\hat{y}^{(t)} = g(vh^{(t)})$, where g can be the **softmax** function.
- ▶ $\text{cost}(y^{(t)}, \hat{y}^{(t)}) = \text{cross_entropy}(y^{(t)}, \hat{y}^{(t)}) = -\sum y^{(t)} \log \hat{y}^{(t)}$
- ▶ $y^{(t)}$ is the **correct** word at time step t , and $\hat{y}^{(t)}$ is the **prediction**.




Recurrent Neurons - Weights

- ▶ Each recurrent neuron has three sets of weights: u , w , and v .
- ▶ u : the weights for the inputs $x^{(t)}$.
- ▶ $x^{(t)}$: is the input at time step t .
- ▶ For example, $x^{(1)}$ could be a one-hot vector corresponding to the first word of a sentence.
- ▶ w : the weights for the hidden state of the previous time step $h^{(t-1)}$.
- ▶ $h^{(t)}$: is the hidden state (memory) at time step t .
 - $h^{(t)} = \tanh(u^T x^{(t)} + w h^{(t-1)})$
 - $h^{(0)}$ is the initial hidden state.



Recurrent Neurons - Weights

- ▶ v : the weights for the hidden state of the current time step $h^{(t)}$.
- ▶ $\hat{y}^{(t)}$ is the output at step t .
- ▶ $\hat{y}^{(t)} = \text{softmax}(vh^{(t)})$
- ▶ For example, if we wanted to predict the next word in a sentence, it would be a vector of probabilities across our vocabulary.





Layers of Recurrent Neurons


- At each time step t , every neuron of a layer receives both the input vector $\mathbf{x}^{(t)}$ and the output vector from the previous time step $\mathbf{h}^{(t-1)}$.

$$\mathbf{h}^{(t)} = \tanh(\mathbf{u}^T \mathbf{x}^{(t)} + \mathbf{w}^T \mathbf{h}^{(t-1)})$$

$$\mathbf{y}^{(t)} = \text{sigmoid}(\mathbf{v}^T \mathbf{h}^{(t)})$$




Time




Deep RNN

- Stacking multiple layers of cells gives you a deep RNN.




Training RNNs

- ▶ To train an RNN, we should unroll it through time and then simply use regular backpropagation.
- ▶ This strategy is called backpropagation through time (BPTT).
- ▶ To train the model using BPTT, we go through the following steps:
 - ▶ 1. Forward pass through the unrolled network (represented by the dashed arrows).
 - ▶ 2. The cost function is $C(\hat{y}^{t_{\min}}, \hat{y}^{t_{\min}+1}, \dots, \hat{y}^{t_{\max}})$, where t_{\min} and t_{\max} are the first and last output time steps, not counting the ignored outputs.



Backpropagation Through Time

- ▶ 3. Propagate backward the gradients of that cost function through the unrolled network (represented by the solid arrows).
- ▶ 4. The model parameters are updated using the gradients computed during BPTT.
- ▶ The gradients flow backward through all the outputs used by the cost function, not just through the final output.
- ▶ For example, in the following figure:
 - The cost function is computed using the last three outputs, $\hat{y}^{(2)}$, $\hat{y}^{(3)}$, and $\hat{y}^{(4)}$.
 - Gradients flow through these three outputs, but not through $\hat{y}^{(0)}$ and $\hat{y}^{(1)}$.



BPTT Step by Step

$$\mathbf{s}^{(t)} = \mathbf{u}^T \mathbf{x}^{(t)} + \mathbf{w} \mathbf{h}^{(t-1)}$$



$$\mathbf{h}^{(t)} = \tanh(\mathbf{s}^{(t)})$$

$$\mathbf{z}^{(t)} = \mathbf{v} \mathbf{h}^{(t)}$$

$$\hat{\mathbf{y}}^{(t)} = \text{softmax}(\mathbf{z}^{(t)})$$

$$J^{(t)} = \text{cross_entropy}(\mathbf{y}^{(t)}, \hat{\mathbf{y}}^{(t)}) = - \sum \mathbf{y}^{(t)} \log \hat{\mathbf{y}}^{(t)}$$

- ▶ We treat the full sequence as **one training example**.
- ▶ The total error E is just the **sum of the errors at each time step**.
- ▶ E.g., $E = J^{(1)} + J^{(2)} + \dots + J^{(t)}$



RNN Design Patterns



Vector-to-Sequence

- ▶ **Vector-to-sequence** network: takes a **single input** at the first time step, and let it output a **sequence**.
- ▶ E.g., the input could be an **image**, and the output could be a **caption for that image**.


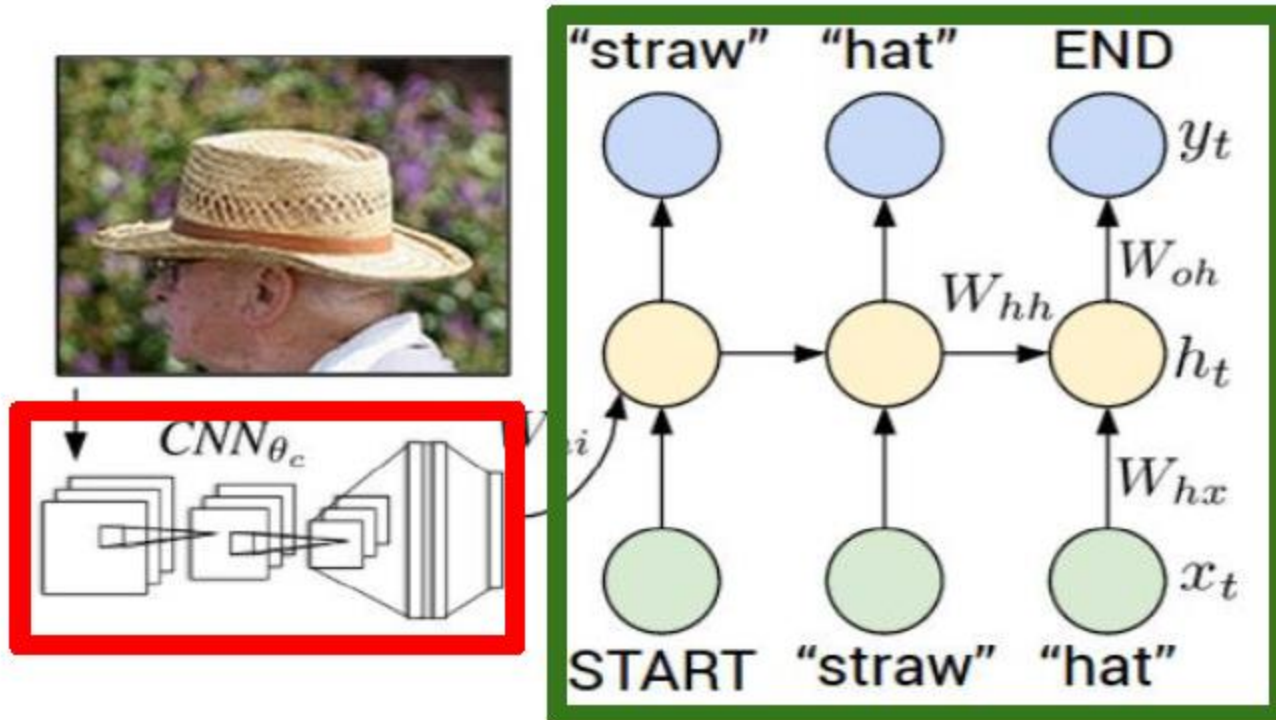


Image Captioning
image \rightarrow sequence of words

Image Captioning

Recurrent Neural Network



Convolutional Neural Network

Image Captioning

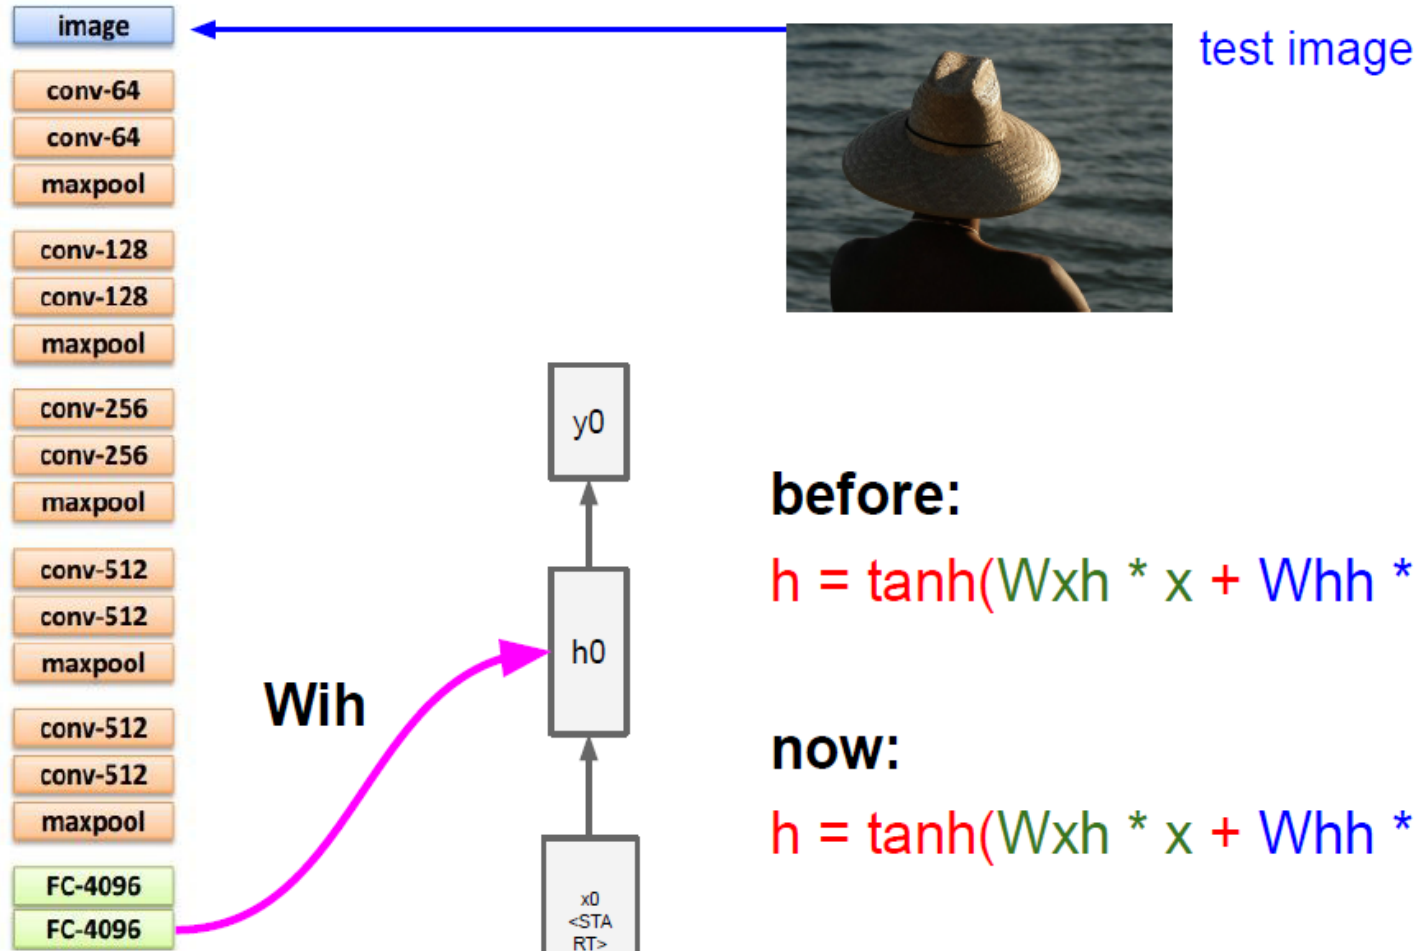
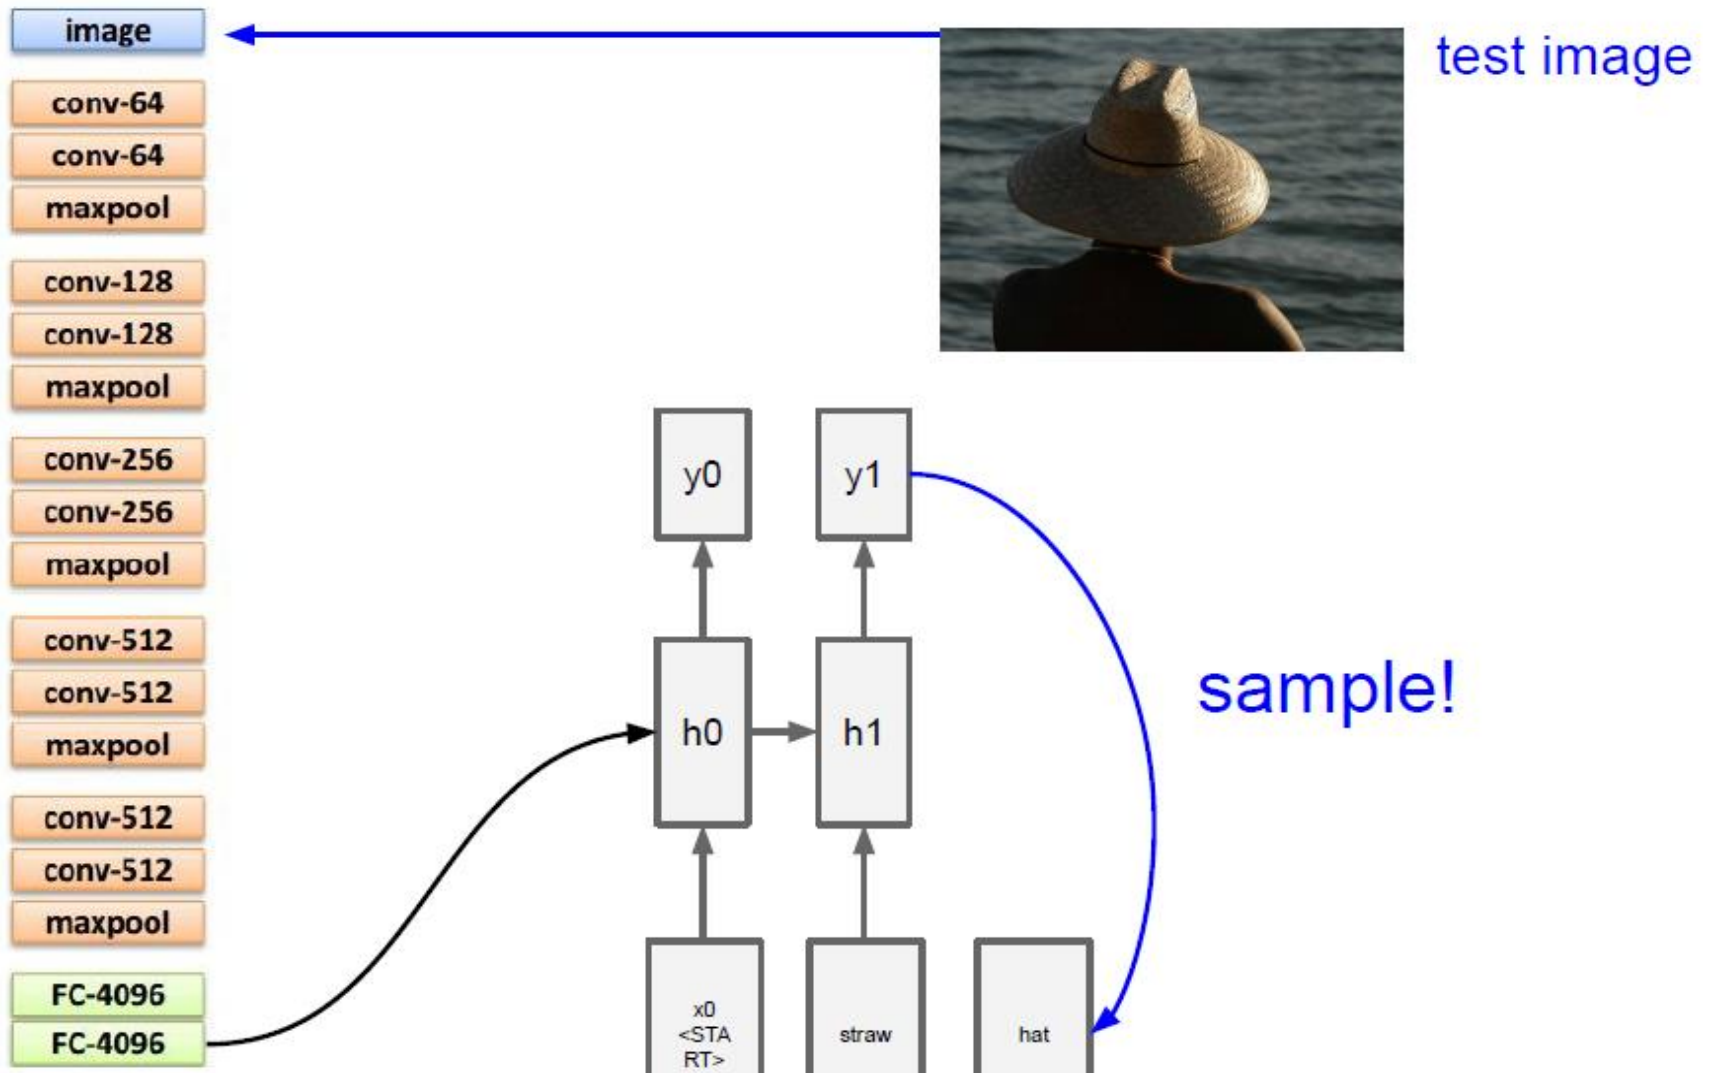



Image Captioning



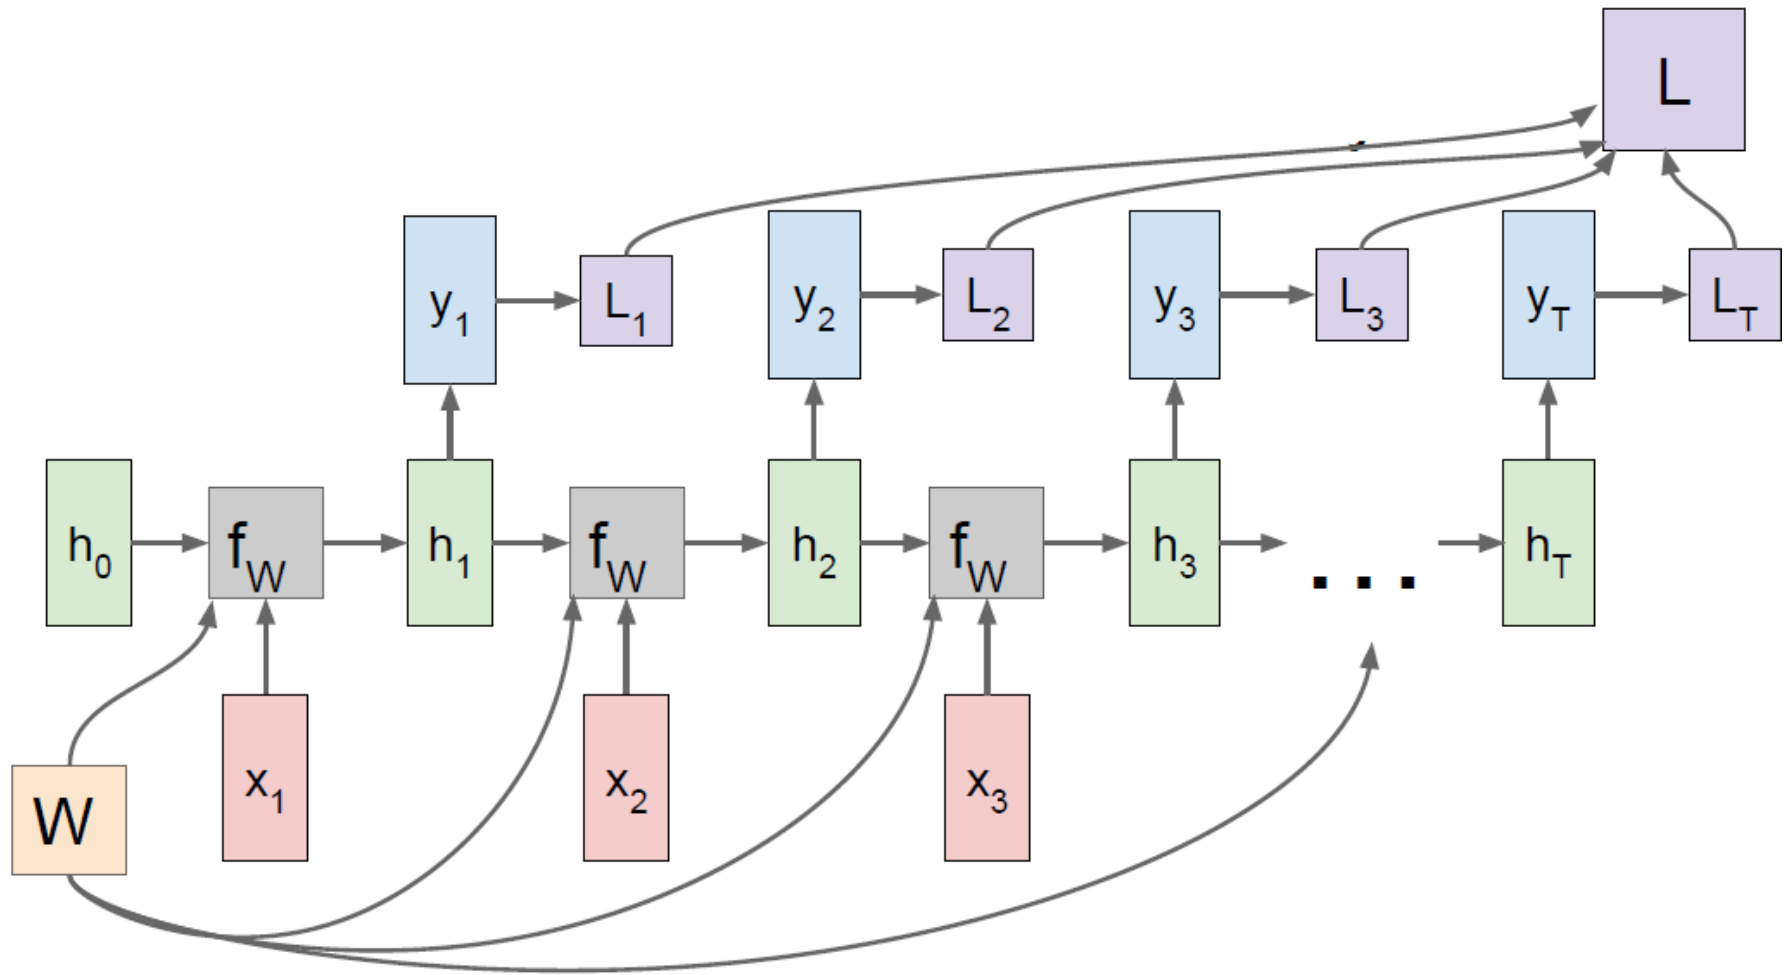
Sequence-to-Vector

- ▶ **Sequence-to-vector** network: takes a **sequence of inputs**, and ignore all outputs except for the **last one**.
- ▶ E.g., you could feed the network a **sequence of words** corresponding to a movie review, and the network would output a **sentiment score**.




Sentiment Classification
sequence of words -> sentiment

Many to Many



Sequence-to-Sequence


- ▶ **Sequence-to-sequence** network: takes a **sequence of inputs** and produce a **sequence of outputs**.
- ▶ Useful for **predicting time series such as stock prices**: you feed it the prices over the last N days, and it must output the prices shifted by one day into the future.
- ▶ Here, both input sequences and output sequences have the **same length**.



Video classification on frame level

Encoder-Decoder

- ▶ **Encoder-decoder** network: a **sequence-to-vector** network (**encoder**), followed by a **vector-to-sequence** network (**decoder**).
- ▶ E.g., **translating** a sentence from one language to another.
- ▶ You would feed the network **a sentence in one language**, the encoder would convert this sentence into a **single vector representation**, and then the decoder would decode this vector into a sentence in another language.

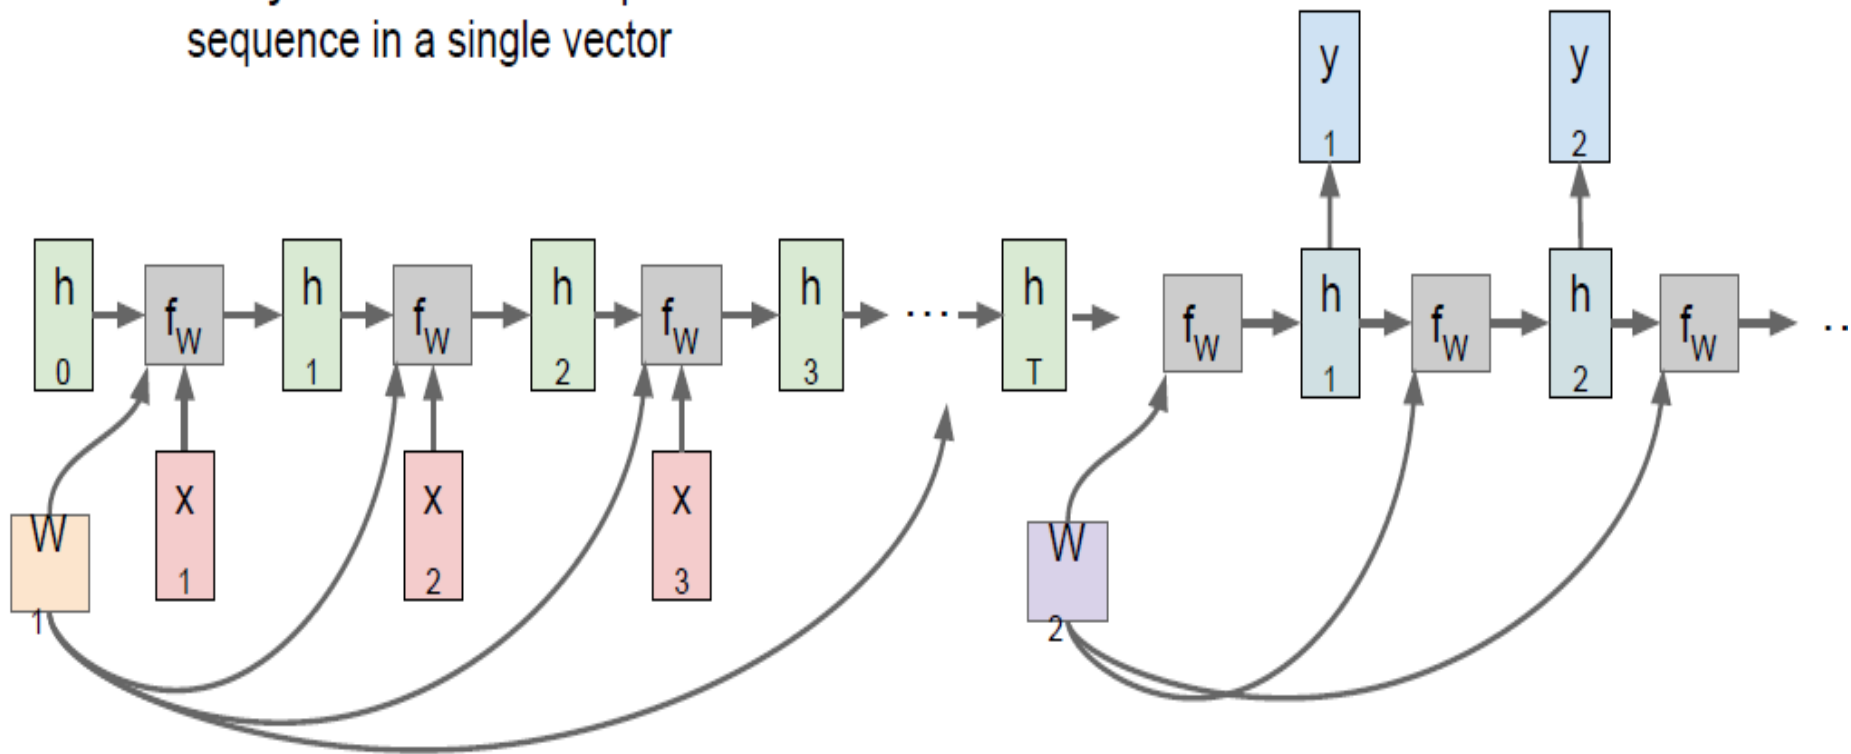


Machine Translation
seq of words \rightarrow seq of words

Many-to-one + one-to-many

Many to one: Encode input sequence in a single vector

One to many: Produce output sequence from single input vector



long short-term memory (LSTM)




RNN Problems


- ▶ Sometimes we only need to look at recent information to perform the present task.
 - E.g., predicting the next word based on the previous ones.
- ▶ In such cases, where the gap between the relevant information and the place that it's needed is small, RNNs can learn to use the past information.
- ▶ But, as that gap grows, RNNs become unable to learn to connect the information.
- ▶ RNNs may suffer from the vanishing/exploding gradients problem.
- ▶ To solve these problem, long short-term memory (LSTM) have been introduced.
- ▶ In LSTM, the network can learn what to store and what to throw away.

RNN Basic Cell vs. LSTM

- ▶ Without looking inside the box, the LSTM cell looks exactly like a basic cell.
- ▶ The repeating module in a standard RNN contains a single layer.




- ▶ The repeating module in an LSTM contains four interacting layers.





LSTM

- In LSTM **state** is split in **two** vectors:
 1. $h^{(t)}$ (**h** stands for **hidden**): the **short-term** state
 2. $c^{(t)}$ (**c** stands for **cell**): the **long-term** state




LSTM

- ▶ The **cell state** (long-term state), the horizontal line on the top of the diagram.
- ▶ The LSTM can **remove/add information** to the **cell state**, regulated by **three gates**.
 - Forget gate, input gate and output gate




LSTM



Step-by-Step LSTM Walk Through

- ▶ **Step one:** decides **what information** we are going to **throw away** from the **cell state**.
- ▶ This decision is made by a **sigmoid layer**, called the **forget gate** layer.
- ▶ It looks at $\mathbf{h}^{(t-1)}$ and $\mathbf{x}^{(t)}$, and outputs a number between 0 and 1 for each number in the cell state $\mathbf{c}^{(t-1)}$.
 - 1 represents **completely keep this**, and 0 represents **completely get rid of this**.


$$\mathbf{f}^{(t)} = \sigma(\mathbf{u}_f^T \mathbf{x}^{(t)} + \mathbf{w}_f \mathbf{h}^{(t-1)})$$



Step-by-Step LSTM Walk Through

- ▶ **Second step:** decides **what new information** we are going to **store** in the **cell state**. This has two parts:
 - ▶ 1. A **sigmoid layer**, called the **input gate layer**, decides **which values** we will update.
 - ▶ 2. A **tanh layer** creates a vector of **new candidate values** that could be added to the state.

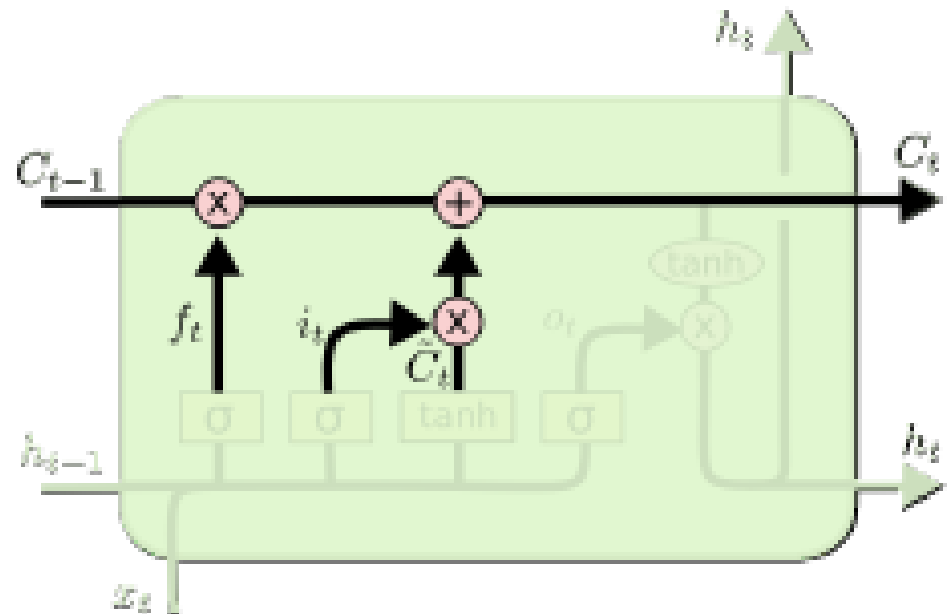
$$i_t^{(t)} = \sigma(\mathbf{u}_i^T \mathbf{x}^{(t)} + \mathbf{w}_i \mathbf{h}^{(t-1)})$$
$$\tilde{c}_t^{(t)} = \tanh(\mathbf{u}_{\tilde{c}}^T \mathbf{x}^{(t)} + \mathbf{w}_{\tilde{c}} \mathbf{h}^{(t-1)})$$



Step-by-Step LSTM Walk Through

- ▶ **Third step:** updates the old cell state $c^{(t-1)}$, into the new cell state $c^{(t)}$.
- ▶ We multiply the old state by $f^{(t)}$, forgetting the things we decided to forget earlier.
- ▶ Then we add it $i^{(t)} \otimes \tilde{c}^{(t)}$.
- ▶ This is the new candidate values, scaled by how much we decided to update each state value.


$$c^{(t)} = f^{(t)} \otimes c^{(t-1)} + i^{(t)} \otimes \tilde{c}^{(t)}$$



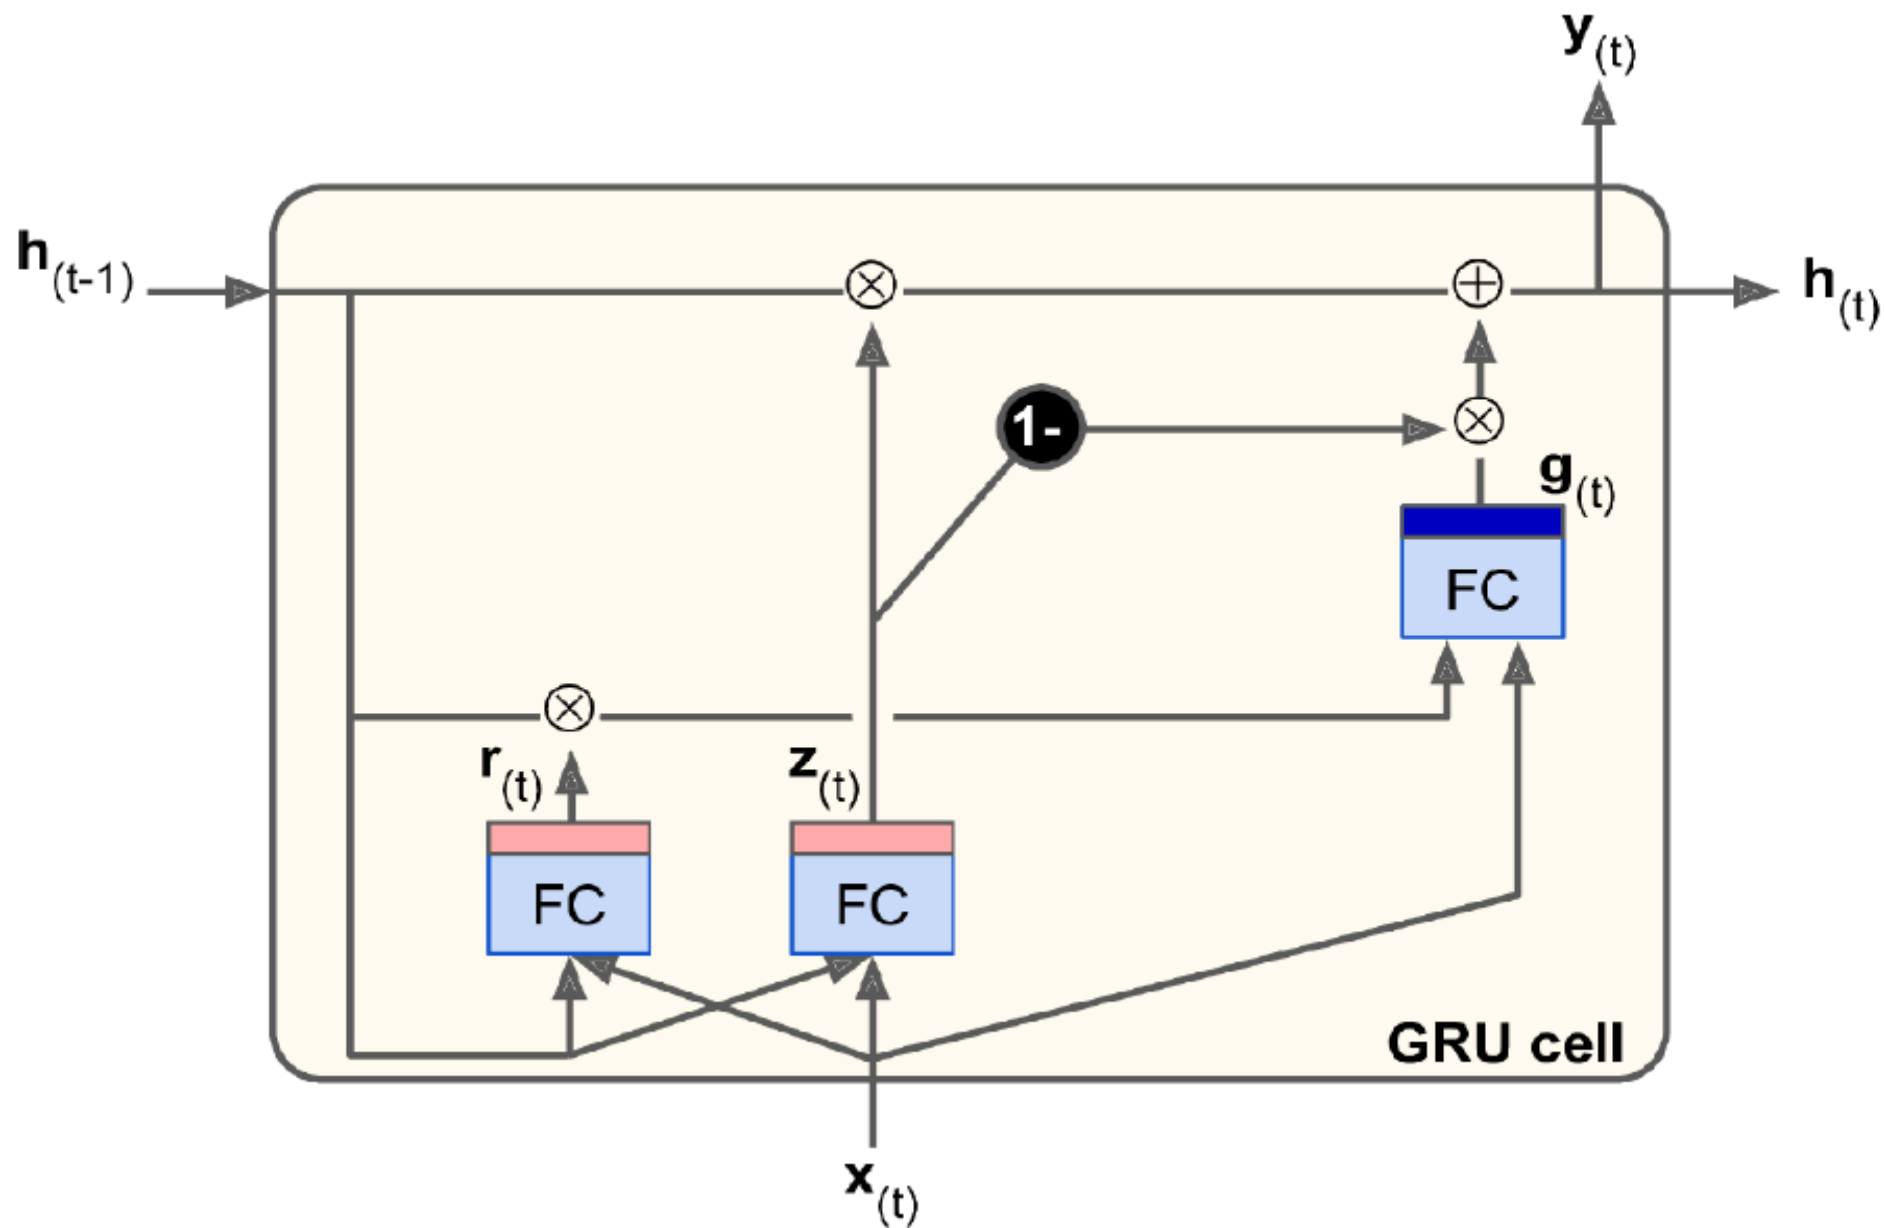
Step-by-Step LSTM Walk Through

- ▶ **Fourth step:** decides about the **output**.
- ▶ First, runs a **sigmoid layer** that decides **what parts of the cell state** we are going to **output**.
- ▶ Then, puts the cell state through **tanh** and multiplies it by the output of the **sigmoid gate (output gate)**, so that it **only outputs the parts it decided to**.

$$o^{(t)} = \sigma(u_o^T x^{(t)} + w_o h^{(t-1)})$$
$$\hat{y}^{(t)} = h^{(t)} = o^{(t)} \otimes \tanh(c^{(t)})$$



Gated Recurrent Unit (GRU)



Gated Recurrent Unit (GRU)

$$\mathbf{z}_{(t)} = \sigma\left(\mathbf{W}_{xz}^\top \mathbf{x}_{(t)} + \mathbf{W}_{hz}^\top \mathbf{h}_{(t-1)} + \mathbf{b}_z\right)$$

$$\mathbf{r}_{(t)} = \sigma\left(\mathbf{W}_{xr}^\top \mathbf{x}_{(t)} + \mathbf{W}_{hr}^\top \mathbf{h}_{(t-1)} + \mathbf{b}_r\right)$$

$$\mathbf{g}_{(t)} = \tanh\left(\mathbf{W}_{xg}^\top \mathbf{x}_{(t)} + \mathbf{W}_{hg}^\top \left(\mathbf{r}_{(t)} \otimes \mathbf{h}_{(t-1)}\right) + \mathbf{b}_g\right)$$

$$\mathbf{h}_{(t)} = \mathbf{z}_{(t)} \otimes \mathbf{h}_{(t-1)} + \left(1 - \mathbf{z}_{(t)}\right) \otimes \mathbf{g}_{(t)}$$

WaveNet- Use 1D convolutional layers

